

Foundations of High Performance Computing

Prelude

How to measure the time consumed by your application

Defining performance in HPC

In you professional life you will have to assess *timing* quite often.

In several cases, the point will be to measure :

- the ***time-to-solution*** that a given code takes to provide a correct result (in general it will be a function of the problem's dimension and of the scalability properties of the code)
- the ***time spent in a code's segment*** if you are profiling it

“Performance” is often related to time a code takes to provide the answer, or for a code's segment, to run.

It may also be translated to the amount of operations (Flops, lops, Memory Accesses, cache misses, ...) that have been performed – that relates, although not linearly, to the time it takes to perform them.

NOTE: in this new era, “performance” may refer to other quantities, for instance to the *energy* it takes to arrive at the problem's solution.

When “**performance**” is related to ***time***, you’re left with the problem of **measuring** the time on the machine you are running on.

So:

- what time do you actually need (or, you have access to)?
- what is the most correct to be used, if many are available ?

What is “time” ?

Basically, you have access to 3 different types of “time”.

1. “wall-clock” time

Basically the same time you can get from the wall-clock in this room.

It is a measure of the “absolute time”.

In POSIX systems, it is the amount of the number of seconds elapsed since the start of the Unix epoch at 1 January 1970 00:00:00 UT.

2. “system-time”

The amount of time that the whole system spent executing your code. It may include I/O, system calls, etc.

3. “process user-time”

The amount of time spent by CPU executing your code’s instructions, strictly speaking.

Where you measure the “time” ?

You can measure all the quoted times :

- ***Outside*** your code
 - you measure the whole code execution
 - you ask the OS to measure the time your code took to execute time:
 - using the `time` command (see `man time`)
 - using `perf` profiler
 - .. discover other ways on your system
- ***Inside*** your code
 - you can measure separate code's section
 - you access system functions to access system's counter

Where you measure the “time” ?

You can measure all the quoted times :

- ***Outside*** your code
 - you measure the whole code execution
 - you ask the OS to measure the time your code took to execute time:
 - using the `time` command (see `man time`)
 - using `perf` profiler
 - .. discover other ways on your system
- ***Inside*** your code
 - you can measure separate code's section
 - you access system functions to access system's counter

That is the focus for the next few slides

How do you measure “time” ?

Baseline: you call the correct system function right before and after the code snippet you’re interested in, and calculate the difference (yes, you’re including the time function’s overhead).

- **Gettimeofday(...)** returns the wall-clock time with μ s precision

Data are given in a `timeval` structure:

```
struct timeval {
    time_t      tv_sec;    /* seconds */
    suseconds_t tv_usec;   /* microseconds */
};
```

- **clock_t clock()** returns the **user-time + system-time** with μ s precision.

Results must be divided by `CLOCKS_PER_SEC`

How do you measure “time” ?

- `int clock_gettime(clockid_t clk_id, struct timespec ..)`

CLOCK_REAL_TIME system-wide realtime clock;

CLOCK_MONOTONIC monotonic time

CLOCK_PROCESS_CPUTIME High-resolution **per-process** timing

CLOCK_THREAD_CPUTIME_ID high-precision **per-thread** timing

Resolution is **1 ns**

```
struct timespec {  
    time_t    tv_sec;           /* seconds */  
    long      tv_nsec;         /* nanoseconds */  
};
```



32-bits quantity: wraps at 136,19 years

How do you measure “time” ?

- **int getrusage(int who, struct rusage *usage)**


RUSAGE_SELF process + all threads

RUSAGE_CHILDREN all the children hierarchy

RUSAGE_THREAD calling thread

Resolution: 1ns

```
struct rusage {
    struct timeval ru_utime; /* user CPU time used */
    struct timeval ru_stime; /* system CPU time used */
    long    ru_maxrss;      /* maximum resident set size */
    long    ru_ixrss;       /* integral shared memory size */
    long    ru_idrss;       /* integral unshared data size */
    long    ru_isrss;       /* integral unshared stack size */
    long    ru_minflt;      /* page reclaims (soft page faults) */
    long    ru_majflt;      /* page faults (hard page faults) */
    long    ru_nswap;       /* swaps */
    long    ru_inblock;     /* block input operations */
    long    ru_oublock;     /* block output operations */
    long    ru_msgsnd;      /* IPC messages sent */
    long    ru_msgrcv;      /* IPC messages received */
    long    ru_nsignals;    /* signals received */
    long    ru_nvcsw;       /* voluntary context switches */
    long    ru_nivcsw;      /* involuntary context switches */
};
```



In practice..

A possibility on a POSIX system is:

```
#define CPU_TIME (clock_gettime( CLOCK_PROCESS_CPUTIME_ID, &ts ), \
                  (double)ts.tv_sec + \
                  (double)ts.tv_nsec * 1e-9)

....

Tstart = CPU_TIME ;

// your code segment here

Time = CPU_TIME - Tstart;
```